

Meuse Data Set Interpolation Exercise

Yin Cheng Ng

December 11, 2015

1 Introduction

We explored various interpolation tools to interpolate soil Cadmium concentration as provided in the Meuse data set. The models that were investigated are Ordinary Least Square (OLS) regression, kernel ridge regression (KRR), Gaussian process (GP), multi-task Gaussian process (MGP) and Bayesian neural network (BNN). The models were compared based on 10-fold cross-validated root mean square error (RMSE) and mean absolute error (MAE). MGP with soil Zinc concentration as co-variable has the lowest cross-validated RMSE and MAE of all models. However, kernel ridge regression has the lowest RMSE and MAE among models with no co-variable.

2 Task description and data set

The key task is to build a model that can interpolate the topsoil cadmium concentration at unmeasured locations, given a set of existing measurements.

2.1 Data set

We investigated the Meuse river data set provided in the 'sp' R package. The variables of our key interest are the following:

- **x, y** - The 'Easting' and 'Northing' map coordinates of the measurements of interest. The coordinate system used follows the Dutch Rijksdriehoek system.
- **cadmium** - The topsoil cadmium concentration, measured in ppm. This variable was incremented by 0.2 compared to the raw measurement, as explained in the description of Meuse data set. This is the target interpolation variable. Refer to Figure 1 for spatial distributions.
- **copper, lead, zinc** - The topsoil copper, lead and zinc concentrations (in ppm). These variables are highly correlated to the topsoil cadmium concentration, and were investigated as covariables in MGP. The Pearson correlation coefficients of the log-concentrations with respect to cadmium log-concentration are tabulated in Table 1.

2.2 Data pre-processing

The 'x' and 'y' coordinate variables in the training data set were standardized. They each have zero sample mean and variance of approximately one. In addition, the top-soil metal concentrations were transformed with base-10 log function for convenience. The original measurements have support from zero to infinity.




Figure 1: Measured log-concentration of topsoil cadmium at various x-y coordinates.

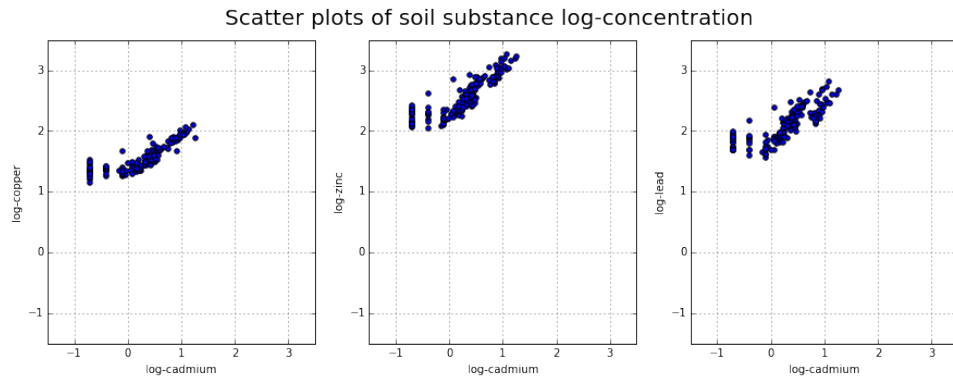


Figure 2: Scatter plots of copper, lead and zinc log-concentrations with respect to cadmium log-concentration.

	Cadmium	Copper	Zinc	Lead
Cadmium	1.0	0.84	0.86	0.81
Copper		1.0	0.90	0.84
Zinc			1.0	0.97
Lead				1.0

Table 1: Pairwise Pearson correlation coefficients of topsoil log-concentration measurements.

3 Models

We investigated the following models for interpolation of topsoil cadmium concentration. Note that target variable in the training data set was centered before model fitting. For predictions, the mean of the initial target variable in the training data set was added to the model output, and the sum was raised to the power of ten to yield outputs comparable to the original measurements in the Meuse data set (i.e., $10^{f_{model} + \mu_{f_{train}}}$ where f is the target variable).

- **OLS Regression (OLS)** - Ordinary least square regression.
- **Kernel Ridge Regression (KRR)** - Kernel ridge regression with Gaussian kernels. Hyperparameters selected through 10-fold cross-validation over a 2-D grid with increment of 0.025 in each dimension.
- **Gaussian Process (GP)** - Gaussian process regression with the Gaussian kernel (also known as the squared-exponential kernel) and Gaussian likelihood. The hyperparameters were selected by maximizing the log-marginal likelihood over the training data set.
- **Cokriging/Multi-task GP (MGP)** - Multi-task GP regression with the Gaussian kernel. We investigated using Zinc, Copper and Lead log-concentrations as covariables. The models are referred to as MGP-Zn, MGP-Cu and MGP-Pb respectively. The hyperparameters were selected by maximizing the log-marginal likelihood. In cross-validation, all measurements of the covariable were included in the training data at all time.
- **Bayesian Neural Network (BNN)** - Bayesian neural network trained with the Probabilistic Backpropagation algorithm proposed in ¹. Architecture with 2 fully-connected hidden layers of 50 units and 60 units were investigated. The models are referred to as BNN-50 and BNN-60 respectively (BNN-50 has 50 hidden units in each hidden layer, BNN-60 has 60 hidden units in each hidden layer).

3.1 Evaluation metrics

The models were compared based on 10-fold cross-validated RMSE and MAE. The results are reported in the results and discussions section.

¹<http://jmlr.org/proceedings/papers/v37/hernandez-lobatoc15.pdf>

Note that GP, MGP and BNN provides predictive distributions as outputs while OLS and KRR only provides point predictions. For GP, MGP and BNN, the model performance was evaluated based on their predictive means.

3.2 Software

Python was the software of choice in the investigation. The OLS and KRR models were fitted with implementations in scikit-learn library ², GP and MGP were fitted using the excellent GPy library ³ and BNNs were fitted using research code from the author of the original paper ⁴.

4 Results and discussions

The cross-validation results of the models are presented in Table 2. Among all the models, MGP-Zn performed the best. However, among the models with no covariable, KRR performed the best while other models, with the exception of OLS, are also competitive.

The relatively large difference between RMSE and MAE among models with no covariable suggests that there may be outliers. This is verified through inspecting the histograms of residuals in Figure 3. There is one residual with magnitude greater than 15 among models with no covariable.

	RMSE	MAE
OLS	3.487	2.078
KRR	2.515	1.511
GP	2.807	1.632
MGP-Zn	1.569	1.056
MGP-Cu	1.902	1.177
MGP-Pb	2.098	1.357
BNN-50	2.714	1.520
BNN-60	2.876	1.587

Table 2: 10-fold cross-validated RMSE and MAE.

²www.scikit-learn.org

³<https://github.com/SheffieldML/GPy>

⁴<https://github.com/HIPS/Probabilistic-Backpropagation>




Figure 3: Histogram of cross-validation residuals for all models. The y-axis denotes the frequency count.